



Roteiro da Aula

1. Correlação vs Regressão Linear
2. Regressão Linear: Conceitos Iniciais
3. Entendendo a regressão linear na prática
4. Como encontrar a reta de regressão ideal? (Minimizando o erro)
5. Exemplo prático: Peso e Altura
 - Conhecendo as bibliotecas: `sklearn` e `statsmodels`
 - Modelo com a presença de outliers
6. Características do Erro
7. Métricas para Análise dos Erros

1. Correlação vs Regressão Linear

Correlação	Regressão
Mede o grau de relação entre duas variáveis	Uma variável afeta a outra
Grau de interrelação	Baseada em causalidade (\sim Relação $\sim \rightarrow$ Causa e efeito)
$\rho(x,y) = \rho(y,x)$	Unidirecional
Único ponto	Linha

\rightarrow Correlação não implica causalidade!

```
In [ ]: y_est = beta_0 + beta_1 * x
        y_verdadeiro = y_est + erro
```

```
In [6]: # Importação das bibliotecas
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

plt.style.use('ggplot') # tema
```

2. Regressão Linear: Conceitos iniciais

```
In [7]: x = np.arange(6)  
y = x
```

```
In [8]: x
```


```
Out[8]: array([0, 1, 2, 3, 4, 5])
```

```
In [9]: y
```

```
Out[9]: array([0, 1, 2, 3, 4, 5])
```

```
In [10]: # Visualização  
plt.scatter(x, y, s=100)
```

```
Out[10]: <matplotlib.collections.PathCollection at 0x22cbe75ec50>
```



Equação da reta de regressão

$$y = b_0 + b_1 * x$$

```
In [11]: # Coeficiente x EQM  
b0 = 0  
  
plt.figure(figsize=(20,10))  
plt.scatter(x, y, s=100)  
  
coeficiente_erro = []  
  
for b1 in range(-3, 6):
```

```

y_est = b0 + b1 * x
plt.plot(x, y_est, label='$\\beta_1 = $' + str(b1))


erro = sum((y_est - y)**2) / 6 # Erro médio quadrático

coeficiente_erro.append([b1, erro])

plt.legend(fontsize=14)

```

Out[11]: <matplotlib.legend.Legend at 0x22cc08a0c40>




In [12]: coeficiente_erro

```
Out[12]: [[-3, 146.66666666666666],
          [-2, 82.5],
          [-1, 36.666666666666664],
          [0, 9.166666666666666],
          [1, 0.0],
          [2, 9.166666666666666],
          [3, 36.666666666666664],
          [4, 82.5],
          [5, 146.66666666666666]]
```

No nosso caso, $y = x$, então, $\beta_0 = 0$, $\beta_1 = 1$.

```
In [13]: plt.plot(coeficiente_erro)
```

```
Out[13]: [<matplotlib.lines.Line2D at 0x22cc0e36fe0>,
          <matplotlib.lines.Line2D at 0x22cc0e36f80>]
```



```
In [14]: coef_erro = np.array(coeficiente_erro)
```

```
In [15]: coef_erro[:,0]
```

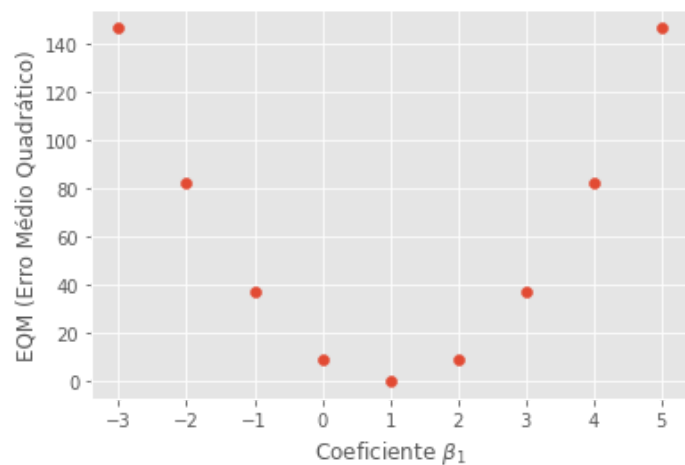
```
Out[15]: array([-3., -2., -1.,  0.,  1.,  2.,  3.,  4.,  5.])
```

```
In [16]: coef_erro[:,1]
```

```
Out[16]: array([[146.6666667,  82.5,  36.6666667,  9.1666667,
                  0.,  9.1666667,  36.6666667,  82.5,
                  146.6666667]])
```

```
In [17]: plt.scatter(coef_erro[:,0], coef_erro[:,1])
plt.xlabel('Coeficiente  $\beta_1$ ')
plt.ylabel('EQM (Erro Médio Quadrático)')
```

```
Out[17]: Text(0, 0.5, 'EQM (Erro Médio Quadrático)')
```



Calculando os coeficientes β_0 e β_1 por meio das equações

$$\beta_1 = \frac{\sum_{i=1}^N (x_i - \bar{x}) \cdot (y_i - \bar{y})}{\sum_{i=1}^N (x_i - \bar{x})^2} \quad \beta_0 = \bar{y} - \beta_1 \bar{x}$$

In [18]:

```
x
```

Out[18]:

```
array([0, 1, 2, 3, 4, 5])
```

In [19]:

```
y
```

Out[19]:

```
array([0, 1, 2, 3, 4, 5])
```

In [20]:

```
x_mean = x.mean()
```

```
y_mean = y.mean()
```

```
x_mean, y_mean
```

Out[20]:

```
(2.5, 2.5)
```

In [21]:

```
x - x_mean
```

Out[21]:

```
array([-2.5, -1.5, -0.5,  0.5,  1.5,  2.5])
```

In [22]:

```
y - y_mean
```

Out[22]:

```
array([-2.5, -1.5, -0.5,  0.5,  1.5,  2.5])
```

In [23]:

```
(y - y_mean) * (y - y_mean)
```

Out[23]:

```
array([6.25, 2.25, 0.25, 0.25, 2.25, 6.25])
```

```
In [24]: sum((y - y_mean) * (y - y_mean))
```

```
Out[24]: 17.5
```

```
In [25]: sum(x - x_mean)
```

```
Out[25]: 0.0
```

```
In [26]: b1 = sum( (x - x_mean) * (y - y_mean) ) / sum( (x - x_mean)**2 )  
b1
```


```
Out[26]: 1.0
```

```
In [27]: b0 = y_mean - (b1 * x_mean)  
b0
```

```
Out[27]: 0.0
```

```
In [28]: plt.scatter(x, y, s=100)  
#plt.plot(x, b0 + b1 * x, color='blue')
```

```
Out[28]: <matplotlib.collections.PathCollection at 0x22cc11c9f30>
```



Regressão Linear Simples: Altura \rightarrow Peso

```
In [30]: df = pd.read_csv('../weight-height.csv')
```

```
In [31]: df.head()
```

Out[31]:

	Gender	Height	Weight
0	Male	73.847017	241.893563
1	Male	68.781904	162.310473
2	Male	74.110105	212.740856
3	Male	71.730978	220.042470
4	Male	69.881796	206.349801

In [32]: `df.describe().T`

Out[32]:

	count	mean	std	min	25%	50%	75%	max
Height	10000.0	66.367560	3.847528	54.263133	63.505620	66.318070	69.174262	78.998742
Weight	10000.0	161.440357	32.108439	64.700127	135.818051	161.212928	187.169525	269.989699

Conversão das unidades

In [33]: `df['altura'] = df['Height'] * 2.54 # pol => cm`

In [34]: `df['peso'] = df['Weight'] * 0.453592 # libras => kg`

In [35]: `df.head()`


Out[35]:

	Gender	Height	Weight	altura	peso
0	Male	73.847017	241.893563	187.571423	109.720985
1	Male	68.781904	162.310473	174.706036	73.622732
2	Male	74.110105	212.740856	188.239668	96.497550
3	Male	71.730978	220.042470	182.196685	99.809504
4	Male	69.881796	206.349801	177.499761	93.598619

Visualização gráfica

In [36]: `sns.scatterplot(data=df, x='altura', y='peso')`

Out[36]: `<AxesSubplot:xlabel='altura', ylabel='peso'>`



Calculando a correlação entre as features

```
In [37]: from scipy.stats import pearsonr
```

```
In [38]: pearsonr(df['altura'], df['peso'])
```

```
Out[38]: (0.9247562987409151, 0.0)
```

- Correlação muito forte: [-1, -0.8] or [0.8, 1]

$$\beta_1 = \frac{\sum_{i=1}^N (x_i - \bar{x}) \cdot (y_i - \bar{y})}{\sum_{i=1}^N (x_i - \bar{x})^2} \quad \beta_0 = \bar{y} - \beta_1 \bar{x}$$

```
In [39]: x = df['altura']
```

```
In [40]: y = df['peso']
```

```
In [41]: x_mean = x.mean()
         y_mean = y.mean()
```

```
x_mean, y_mean
```

```
Out[41]: (168.57360177724598, 73.22805433651739)
```

```
In [42]: b1 = sum( (x - x_mean) * (y - y_mean) ) / sum( (x - x_mean)**2 )
         b1
```

```
Out[42]: 1.3781495809287967
```

```
In [43]: b0 = y_mean - (b1 * x_mean)
```




```
b0
```

```
Out[43]: -159.091584308452
```

```
In [44]: plt.scatter(df['altura'], df['peso'])  
plt.plot(x, b0 + b1 * x, color='blue', linewidth=4)
```

```
Out[44]: [<matplotlib.lines.Line2D at 0x22cc1c1b400>]
```



```
In [45]: y_est = b0 + b1 * x
```

```
In [46]: y
```

```
Out[46]: 0      109.720985  
1       73.622732  
2       96.497550  
3       99.809504  
4       93.598619  
...  
9995    62.041159  
9996    77.504315  
9997    58.275377  
9998    74.322166  
9999    51.550324  
Name: peso, Length: 10000, dtype: float64
```

```
In [47]: erro = y - y_est  
erro
```

```
Out[47]: 0      10.311091
         1      -8.056735
         2      -3.833285
         3       7.806803
         4       8.068981
         ...
        9995 -10.504621
        9996  1.827329
        9997 -6.202942
        9998 -8.240614
        9999 -6.193921
Length: 10000, dtype: float64
```

Scikit Learning

```
In [49]: from sklearn.linear_model import LinearRegression
```

```
In [50]: lr = LinearRegression()
```

```
In [51]: lr
```

```
Out[51]: ▼ LinearRegression
          LinearRegression()
```

```
In [52]: df['altura']
```

```
Out[52]: 0      187.571423
         1      174.706036
         2      188.239668
         3      182.196685
         4      177.499761
         ...
        9995  168.078536
        9996  170.350573
        9997  162.224700
        9998  175.346978
        9999  157.338385
Name: altura, Length: 10000, dtype: float64
```

```
In [53]: X = df['altura'].values
         y = df['peso'].values
```

```
In [54]: X
```

```
Out[54]: array([187.57142322, 174.70603628, 188.2396677 , ..., 162.22470022,
                175.34697755, 157.33838453])
```

```
In [55]: y
```

```
Out[55]: array([109.72098511,  73.62273185,  96.49755015, ...,  58.2753768 ,
                74.32216565,  51.55032378])
```

Entendendo o método `reshape`

```
In [56]: array = np.ones(25)
          array
```

```
Out[56]: array([1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,  
                1., 1., 1., 1., 1., 1., 1., 1.])
```

```
In [57]: array.reshape(5, 5)
```

```
Out[57]: array([[1., 1., 1., 1., 1.],
        [1., 1., 1., 1., 1.],
        [1., 1., 1., 1., 1.],
        [1., 1., 1., 1., 1.],
        [1., 1., 1., 1., 1.]])
```

```
In [58]: X_reshape = X.reshape(-1, 1)
```

```
In [59]: X_reshape.shape # Vetor coluna
```

```
Out[59]: (10000, 1)
```

```
In [60]: X.shape      # Vetor Linha
```

```
Out[60]: (10000,)
```

```
In [61]: # Treinamento do modelo: calcular os parâmetros do modelo
```

```
lr.fit(X.reshape(-1, 1), y)
```

```
Out[61]: ▼ LinearRegression
LinearRegression()
```

```
In [62]: lr.coef_ # coeficientes
```

```
Out[62]: array([1.37814958])
```

```
In [63]: lr.intercept_ # intercepto
```

```
Out[63]: -159.09158430845105
```

Statsmodels

```
In [67]: import statsmodels.api as sm
```

```
In [68]: x = sm.add_constant(X)

model = sm.OLS(y, x).fit()
```

```
In [69]: model.summary()
```

Out[69]:

OLS Regression Results

Dep. Variable:	y	R-squared:	0.855
Model:	OLS	Adj. R-squared:	0.855
Method:	Least Squares	F-statistic:	5.904e+04
Date:	Tue, 14 Jun 2022	Prob (F-statistic):	0.00
Time:	10:01:33	Log-Likelihood:	-31313.
No. Observations:	10000	AIC:	6.263e+04
Df Residuals:	9998	BIC:	6.265e+04
Df Model:	1		
Covariance Type:	nonrobust		
	coef	std err	t P> t [0.025 0.975]
const	-159.0916	0.958	-166.109 0.000 -160.969 -157.214
x1	1.3781	0.006	242.975 0.000 1.367 1.389
Omnibus:	2.141	Durbin-Watson:	1.677
Prob(Omnibus):	0.343	Jarque-Bera (JB):	2.150
Skew:	0.036	Prob(JB):	0.341
Kurtosis:	2.991	Cond. No.	2.92e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.92e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Influência dos Outliers

```
In [70]: df.head()
```

```
Out[70]:
```


	Gender	Height	Weight	altura	peso
0	Male	73.847017	241.893563	187.571423	109.720985
1	Male	68.781904	162.310473	174.706036	73.622732
2	Male	74.110105	212.740856	188.239668	96.497550
3	Male	71.730978	220.042470	182.196685	99.809504
4	Male	69.881796	206.349801	177.499761	93.598619

Obtendo apenas 100 amostras (como array)

```
In [71]: X = df['altura'].values[:100]  
y = df['peso'].values[:100]
```

```
In [72]: # Visualizando  
plt.scatter(X, y)
```

```
Out[72]: <matplotlib.collections.PathCollection at 0x22cc5ad1f00>
```



Modelo sem a presença dos outliers

```
In [73]: lr = LinearRegression()
```

```
In [74]: lr.fit(X.reshape(-1, 1), y)
```

```
Out[74]: ▼ LinearRegression
LinearRegression()
```

```
In [75]: lr.coef_[0]
```


```
Out[75]: 1.0449527098060838
```

```
In [76]: lr.intercept_
```

```
Out[76]: -98.52075866207376
```

```
In [77]: plt.scatter(X, y)
plt.plot(X, lr.intercept_ + lr.coef_[0] * X, color='blue', linewidth=4)
```

```
Out[77]: <matplotlib.lines.Line2D at 0x22cc5b4ac20>
```



```
In [78]: lr.score(X.reshape(-1, 1), y)
```


```
Out[78]: 0.6214261131940761
```

Modelo com a presença dos outliers

```
In [79]: X2 = np.append(X, [150, 140, 173])
y2 = np.append(y, [200, 100, 250])
```

```
In [80]: plt.scatter(X2, y2)
```

```
Out[80]: <matplotlib.collections.PathCollection at 0x22cc5bb6440>
```



```
In [81]: lr2 = LinearRegression() # Modelo com os outliers
```

```
In [82]: lr2.fit(X2.reshape(-1, 1), y2)
```

```
Out[82]: ▼ LinearRegression
LinearRegression()
```

```
In [83]: lr2.coef_[0]
```

```
Out[83]: 0.1348777878056209
```

```
In [84]: lr2.intercept_
```

```
Out[84]: 64.05035131391783
```

```
In [85]: plt.figure(figsize=(20, 10))
```

```
# plot de dispersão dos dados originais com os 3 outliers
plt.scatter(X2, y2)
```


```
# plot da reta de regressão do modelo treinado com os dados + outliers
```

```
plt.plot(X2, lr2.intercept_ + lr2.coef_[0] * X2, color='blue', linewidth=4, label='Modelo com outliers')
```

```
plt.plot(X2, lr2.intercept_ + lr2.coef_[0] * X2, color='green', linewidth=4, label='Modelo sem outliers')
```

```
plt.legend(fontsize=20)
```


```
Out[85]: <matplotlib.legend.Legend at 0x22cc5bdc520>
```



```
In [86]: # plot de dispersão dos dados originais com os 3 outliers
plt.scatter(X, y)

# plot da reta de regressão do modelo treinado com os dados + outliers
plt.plot(X, lr2.intercept_ + lr2.coef_[0] * X, color='blue', linewidth=4, label='Modelo com outliers')
plt.plot(X, lr.coef_[0] * X, color='green', linewidth=4, label='Modelo sem outliers')
```

```
Out[86]: [<matplotlib.lines.Line2D at 0x22cc5cbe560>]
```

Características dos Resíduos

1. Não devem ser correlacionados

Sem outliers

```
In [87]: y1_est = lr.intercept_ + lr.coef_[0] * X # sem outliers
```


```
In [88]: residuos = y - y1_est
```

```
In [89]: residuos
```

```
Out[89]: array([ 12.23847679, -10.41605551, -1.68324204,  7.94334299,
        6.64052067, -10.93862649, -0.61900139, -6.69814429,
        0.44059742,  1.03745092, -5.80209135,  5.32445947,
        2.42695011,  0.56230009, -0.6995153 , -2.92027033,
       -4.77351824,  7.13994933, -1.95565208,  0.95364938,
       -2.76040881,  3.90538528, -9.38700051,  2.67509001,
       -8.19507602,  4.89992669, -1.19723156, -0.89016356,
       -0.95171279, -13.27246275,  1.44755404, -1.40390926,
       -2.86292813,  7.97381595,  6.45756184, -0.83647932,
        4.7592617 ,  3.55224787, -1.82181775,  1.83036849,
       -3.81602721, -5.06552018,  1.76518136, -3.77840911,
       -5.97519411, -1.85630938,  2.21727739,  1.7942771 ,
        2.0011387 ,  0.97298052,  1.6962389 ,  3.44560591,
       -4.60414973,  2.0976939 , -2.29359128, -3.90511334,
        3.15048432, -2.18109307, -0.59787555,  7.73223693,
       -0.84027929,  3.76451322, -8.82465138, -4.48405706,
       -9.45186959, -4.3218979 ,  1.00529284,  5.48617176,
        0.9219087 ,  4.76525105, -1.65497362,  5.47467014,
       -6.00237573, -6.90475905,  4.14602364,  2.23261243,
       -0.93777685,  6.96836941, 11.09236542, -0.23805277,
        1.45526328,  7.49207415, 10.49006559, -8.25740867,
       -3.66288666, -4.58696462, 16.26641358,  3.25160357,
       -3.63755472, -3.89629064, -0.63676411,  7.47547717,
       -6.36261256,  7.24550945, -4.75298696,  0.8551966 ,
       -4.20124643,  3.09104347,  0.4628156 , -5.21066612])
```

```
In [90]: plt.scatter(y1_est, residuos)
```

```
Out[90]: <matplotlib.collections.PathCollection at 0x22cc5d21960>
```



Com outliers

```
In [91]: y2_est = lr2.intercept_ + lr2.coef_[0] * X2
```


```
In [92]: residuos2 = y2 - y2_est
```

```
In [93]: residuos2
```

```
Out[93]: array([ 2.03714152e+01, -1.39915832e+01,  7.05784888e+00,  1.11848670e+01,
        5.60749229e+00, -1.80483176e+01, -4.18718463e+00, -1.12754887e+01,
       -7.21015847e+00, -1.48482430e+01, -3.79864611e+00,  8.35753983e+00,
       -1.04309416e+01, -1.85473696e+00, -3.20748539e+00, -9.12242185e+00,
        5.68501780e-02, -7.54872974e+00, -3.54747497e+00, -4.57725063e+00,
       -8.33974740e+00,  1.84981107e+00, -1.90484186e+01,  1.39493793e+01,
       -1.38249318e+01, -1.50097935e-01, -4.06141957e+00,  4.82445030e+00,
       -7.67253297e+00, -1.75079233e+01, -2.50659706e+00,  7.16081248e+00,
       -9.01076042e-02,  5.31889545e+00,  4.72043809e+00, -6.69484836e-01,
        1.86395908e+00, -3.69496533e+00, -1.22733751e+01, -1.20958044e+01,
       -9.25832078e+00, -7.64253209e-01, -9.92763090e+00, -1.35761060e+01,
       -1.24900996e+01, -2.37387533e+00, -2.58354939e+00,  6.05690158e+00,
       -6.54483323e-01,  6.84828835e+00, -1.11240805e+01,  3.11010921e+00,
       -8.85108754e+00, -4.83212039e+00, -1.00090846e+01, -1.27565087e+01,
       -1.41636280e+00, -1.15064984e+00,  1.45305297e+00,  7.00334330e+00,
        1.64351428e+00,  1.96914369e+00, -7.57552776e-01, -1.21858226e+01,
       -6.93274643e+00, -1.59940106e+01, -3.74385825e+00,  1.11187853e+01,
       -4.23133536e+00,  1.45509866e+00, -4.05712434e+00,  5.92436203e+00,
       -7.53949408e+00, -1.32233012e+01,  9.18398154e+00,  4.36921296e+00,
        2.48453318e+00,  1.13984833e+01,  2.27120049e+00, -1.01770881e+01,
       -4.90515301e+00,  2.67907029e+00,  1.85747148e+01, -9.47381898e+00,
       -1.37624763e+01, -2.85250690e+00,  7.63842567e+00,  6.81934165e-02,
       -8.45542866e+00,  6.99471001e-01,  2.79995535e+00, -4.29430181e+00,
       -5.64175826e+00, -5.69190909e+00, -1.22081788e+01, -1.11909668e+01,
       -1.02847921e+00, -4.98983340e+00, -8.45393913e+00, -7.33076524e+00,
        1.15717981e+02,  1.70667584e+01,  1.62615791e+02])
```

```
In [94]: plt.scatter(y2_est, residuos2)
```

```
Out[94]: <matplotlib.collections.PathCollection at 0x22cc5d9cc40>
```



2. Ter média zero

Sem outliers

```
In [95]: residuos.mean()
```

```
Out[95]: 3.467448550509289e-14
```

Com outliers

```
In [96]: residuos2.mean()
```

```
Out[96]: 1.2969129545912509e-14
```

4. Distribuição normal dos resíduos

Sem outliers


```
In [97]: from scipy.stats import normaltest
```

```
In [98]: normaltest(residuos)
```

```
Out[98]: NormaltestResult(statistic=0.9490854261141337, pvalue=0.6221695014128703)
```

```
In [99]: sns.histplot(residuos, kde=True)
```


```
Out[99]: <AxesSubplot:ylabel='Count'>
```



Com outliers

```
In [100]: sns.histplot(residuos2, kde=True)
```

Out[100]: <AxesSubplot:ylabel='Count'>



In [101]: `model`

Out[101]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x22cc0cc91e0>

In [102]: `model.resid`


Out[102]: array([10.31109111, -8.05673452, -3.83328469, ..., -6.20294152,
-8.24061367, -6.19392062])

In [103]: `normaltest(model.resid)`

Out[103]: NormaltestResult(statistic=2.1407382479464676, pvalue=0.3428819281169143)

In [104]: `sns.histplot(model.resid)`

Out[104]: <AxesSubplot:ylabel='Count'>



Métricas para Análise dos Erros

$$SSE = \sum_{i=1}^N \epsilon_i^2 = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2 = \frac{1}{N} \sum_{i=1}^N (y_i - (\beta_0 + \beta_1 x_i))^2$$
$$SQR = \sum_{i=1}^N (\hat{y}_i - \overline{y})^2$$
$$SQT = \sum_{i=1}^N (y_i - \overline{y})^2$$

1. R-Quadrado (R^2)

$$R^2 = 1 - \frac{SSE}{SQT} = \frac{SQT - SSE}{SQT} = \frac{SQR}{SQT}$$

```
In [105]: from sklearn.metrics import r2_score
```

```
In [106]: lr.score(X.reshape(-1, 1), y)
```

```
Out[106]: 0.6214261131940761
```

```
In [107]: r2_score(y, y1_est)
```

```
Out[107]: 0.6214261131940761
```

2. MAE

```
In [108]: from sklearn.metrics import mean_absolute_error
```

```
In [109]: mean_absolute_error(y, y1_est)
```

```
Out[109]: 4.260493303495109
```

```
In [110]: mean_absolute_error(y2, y2_est)
```

```
Out[110]: 9.640274384813585
```

3. MSE

```
In [111]: from sklearn.metrics import mean_squared_error
```

```
In [112]: mean_squared_error(y, y1_est)
```

```
Out[112]: 28.728718184843036
```

```
In [113]: mean_squared_error(y2, y2_est)
```

```
Out[113]: 460.66295293941755
```

4. RMSE

```
In [114]: np.sqrt(mean_squared_error(y, y1_est))
```

```
Out[114]: 5.359917740492202
```

```
In [115]: np.sqrt(mean_squared_error(y2, y2_est))
```

```
Out[115]: 21.4630601951217
```

```
In [ ]:
```

5. RMSLE

```
In [116]: from sklearn.metrics import mean_squared_log_error
```

```
In [117]: mean_squared_log_error(y, y1_est)
```

```
Out[117]: 0.003965422358995947
```

```
In [118]: mean_squared_log_error(y2, y2_est)
```

```
Out[118]: 0.027940307257655395
```

```
In [ ]:
```